

# Pet breed detection with ResNet50

## Data import and discovery

### Library and data import

```
In [1]: import tensorflow as tf
import tensorflow_datasets as tfds
from tensorflow import keras
from tensorflow.keras import layers, models
from tensorflow.keras.applications.resnet50 import preprocess_input, ResNet50

import numpy as np
from numpy import round, sqrt, random
import matplotlib.pyplot as plt

dataset, info = tfds.load('oxford_iiit_pet', split=['train[:80%]', 'train[80%:]', 'test'], as_supervised=True, with_info=True)
train_set_raw, valid_set_raw, test_set_raw = dataset

# Create a dictionary of numerical label - breed key-value pairs
labels = info.features['label'].names # Name of the breed
label_dict = {i: breed for i, breed in enumerate(labels)} # Dict comprehension to create key-value pairs of numerical label and breed
print(label_dict)

# Find out about the pixel value range
for image, _ in train_set_raw.take(1):
    print(f"Pixel value range: [{tf.reduce_min(image).numpy()}, {tf.reduce_max(image).numpy()}]")

# Number of classes
num_classes = info.features['label'].num_classes
print(f"There are {num_classes} classes of dogs and cats in this dataset")

{0: 'Abyssinian', 1: 'american_bulldog', 2: 'american_pit_bull_terrier', 3: 'basset_hound', 4: 'beagle', 5: 'Bengal', 6: 'Birman', 7: 'Bombay', 8: 'boxer', 9: 'British_Shorthair', 10: 'chihuahua', 11: 'Egyptian_Mau', 12: 'english_cocker_spaniel', 13: 'english_setter', 14: 'german_shorthaired', 15: 'great_pyrenees', 16: 'havanese', 17: 'japanese_chin', 18: 'keeshond', 19: 'leonberger', 20: 'Maine_Coon', 21: 'miniature_pinscher', 22: 'newfoundland', 23: 'Persian', 24: 'pomeranian', 25: 'pug', 26: 'Ragdoll', 27: 'Russian_Blue', 28: 'saint_bernard', 29: 'samoyed', 30: 'scottish_terrier', 31: 'shiba_inu', 32: 'Siamese', 33: 'Sphynx', 34: 'staffordshire_bull_terrier', 35: 'wheaton_terrier', 36: 'yorkshire_terrier'}
Pixel value range: [0, 255]
There are 37 classes of dogs and cats in this dataset
```

### Input data visualization

In this section, I will visualize the images and labels from the dataset

```
In [2]: num_samples = 9
num_rows = int(round(sqrt(num_samples))); num_cols = int(num_samples/num_rows)
index = 1
plt.figure(figsize=(num_rows*3, num_cols*3))

# Display some images from the raw train set
for image, label in train_set_raw.take(num_samples):
    # print(image)
    plt.subplot(num_rows, num_cols, index)
    plt.imshow(image.numpy().astype("uint8"))
    plt.title(f"Label: {label_dict[label.numpy()]}")
    plt.axis("off")
    index += 1
plt.show()
plt.tight_layout
```

Label: Sphynx




Label: english\_cocker\_spaniel



Label: British\_Shorthair



Label: Siamese



Label: Sphynx



Label: american\_pit\_bull\_terrier



Label: newfoundland



Label: newfoundland



Label: yorkshire\_terrier



```
Out[2]: <function matplotlib.pyplot.tight_layout(*, pad: 'float' = 1.08, h_pad: 'float | None' = None, w_pad: 'float | None' = None, rect: 'tuple[float, float, float, float] | None' = None) -> 'None'>
```

## Examine the image dimensions

As can be seen from the visualization above, each image has a unique size. I have created the function below to determine the minimum height and width resolution of the images in the train, validation, and test set

```
In [3]: def find_min_resolution(dataset):

    min_height = float('inf')
    min_width = float('inf')
    for image, _ in dataset:
        height,width,_ = image.shape

        min_height = min(min_height,height)
        min_width = min(min_width,width)

    print(f"The minimum height resolution: {min_height}\n"
          f"The minimum width resolution: {min_width}")

    find_min_resolution(train_set_raw)
    find_min_resolution(valid_set_raw)
    find_min_resolution(test_set_raw)
```

```
The minimum height resolution: 108
The minimum width resolution: 114
The minimum height resolution: 112
The minimum width resolution: 150
The minimum height resolution: 103
The minimum width resolution: 137
```

## Image resizing

So the minimum height resolution is 103 and the minimum width resolution is 114. I will resize the images to have dimension of (224,224) for the pretrained pre-trained ResNet50 network. Normally, I would include the image preprocessing step inside the final model. However, given the variety of image dimensions, I will have to resize the images first

```
In [4]: def preprocess_image(image, target_size = (96, 96),
        display=False,
        pad=True):
    '''This function resize an image while keeping the aspect ratio such that the shorter side is 96 pixels (by

    if pad:      # Resize with padding (shrink the image while preserving aspect ratio and fill void with black)
        image = tf.image.resize_with_pad(image, target_size[0], target_size[1])
    else:        # Resize without padding - stretch or shrink the image to the desired target size
        image = tf.image.resize(image, target_size)

    return image

def preprocess_dataset(dataset, target_size = (96,96), display=False, pad=True):
    '''This function applies the preprocess_image() on the images in a dataset to resize the iamges to the targ
    return dataset.map(lambda image, label: (preprocess_image(image, target_size, display, pad), label))

# The desired size for the processed images
TARGET_SIZE = (224,224)
train_set_processed = preprocess_dataset(train_set_raw, target_size=TARGET_SIZE)
valid_set_processed = preprocess_dataset(valid_set_raw, target_size=TARGET_SIZE)
test_set_processed = preprocess_dataset(test_set_raw, target_size=TARGET_SIZE)

train_set_processed_display = preprocess_dataset(train_set_raw, target_size=TARGET_SIZE, display=True)

find_min_resolution(train_set_processed_display)
```

The minimum height resolution: 224

The minimum width resolution: 224

```
In [5]: num_samples = 9
num_rows = int(round(sqrt(num_samples))); num_cols = int(num_samples/num_rows)
index = 1
plt.figure(figsize=(num_rows*3, num_cols*3))

# Display some images from the raw train set
for image, label in train_set_processed_display.take(num_samples):
    ## Examine the range of pixel value
    # print("Pixel value range:", tf.reduce_min(image).numpy(), "to", tf.reduce_max(image).numpy())

    plt.subplot(num_rows, num_cols, index)
    plt.imshow(image.numpy().astype("uint8"))
    plt.title(f"Label: {label_dict[label.numpy()]}")
    plt.axis("off")
    index += 1
plt.show()
plt.tight_layout
```


Label: Sphynx




Label: english\_cocker\_spaniel




Label: British\_Shorthair



Label: Siamese




Label: Sphynx




Label: american\_pit\_bull\_terrier




Label: newfoundland



Label: newfoundland



Label: yorkshire\_terrier



```
Out[5]: <function matplotlib.pyplot.tight_layout(*, pad: 'float' = 1.08, h_pad: 'float | None' = None, w_pad: 'float | None' = None, rect: 'tuple[float, float, float, float] | None' = None) -> 'None'>
```

## Dataset => Array Block

In this code block, I define a function to extract the image and label data from any of the three datasets. This make it easier to work with the data in the form of numpy arrays instead of PrefetchDataset objects when importing the dataset from tensorflow.

Since the dataset has images of different dimensions, this function will also resize the images to a

```
In [6]: def get_npararray_dataset(dataset):
        image_list = []
        label_list = []

        for image,label in dataset:
            image_list.append(image.numpy())
            label_list.append(label.numpy())

        image_list_npararray = np.array(image_list)
        label_list_npararray = np.array(label_list)

        return image_list_npararray, label_list_npararray
```

```
In [7]: train_image_array, train_label_array = get_npararray_dataset(train_set_processed)
        valid_image_array, valid_label_array = get_npararray_dataset(valid_set_processed)
        test_image_array, test_label_array = get_npararray_dataset(test_set_processed)

        print(train_image_array.shape, train_label_array.shape)
        print(valid_image_array.shape, valid_label_array.shape)
        print(test_image_array.shape, test_label_array.shape)

        num_train = len(train_label_array)
        num_valid = len(valid_label_array)
        num_test = len(test_image_array)
```

```
(2944, 224, 224, 3) (2944,)
(736, 224, 224, 3) (736,)
(3669, 224, 224, 3) (3669,)
```

# ResNet50-based model

## Base model construction

```
In [8]: base_model = ResNet50(weights='imagenet',
    #   input_shape=TARGET_SIZE+(3,),
    include_top=False)

# Freeze the weights of the model
base_model.trainable = False

# # Investigate the structure of the base model and make sure that the weights are frozen
# base_model.summary()
```

## Full model architecture

```
In [9]: model_resnet = models.Sequential([
    layers.Input(shape=TARGET_SIZE + (3,)),
    layers.Lambda(preprocess_input),
    base_model,
    layers.GlobalAveragePooling2D(),
    layers.BatchNormalization(),
    layers.Dropout(0.5),
    layers.Dense(num_classes, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(num_classes, activation='softmax') # Output layer for pet breeds
])

initial_weights = model_resnet.get_weights()

model_resnet.summary()
```

WARNING:tensorflow:From c:\mnguyen\TME\_6015\.venv\Lib\site-packages\keras\src\backend\tensorflow\core.py:204: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From c:\mnguyen\TME\_6015\.venv\Lib\site-packages\keras\src\backend\tensorflow\core.py:204: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

Model: "sequential"

| Layer (type)                                      | Output Shape        | Param #    |
|---|---------------------|------------|
| lambda (Lambda)                                   | (None, 224, 224, 3) | 0          |
| resnet50 (Functional)                             | (None, 7, 7, 2048)  | 23,587,712 |
| global_average_pooling2d (GlobalAveragePooling2D) | (None, 2048)        | 0          |
| batch_normalization (BatchNormalization)          | (None, 2048)        | 8,192      |
| dropout (Dropout)                                 | (None, 2048)        | 0          |
| dense (Dense)                                     | (None, 37)          | 75,813     |
| dropout_1 (Dropout)                               | (None, 37)          | 0          |
| dense_1 (Dense)                                   | (None, 37)          | 1,406      |

Total params: 23,673,123 (90.31 MB)

Trainable params: 81,315 (317.64 KB)

Non-trainable params: 23,591,808 (90.00 MB)

## Model Training

```
In [10]: model_resnet.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.005),
    loss=tf.keras.losses.SparseCategoricalCrossentropy,
    metrics=["accuracy"])
loss0, acc0 = model_resnet.evaluate(test_image_array, test_label_array)

print("initial loss: {:.2f}".format(loss0))
print("initial accuracy: {:.2f}".format(acc0))
```

115/115 ————— 41s 346ms/step - accuracy: 0.0306 - loss: 3.8299  
initial loss: 3.80  
initial accuracy: 0.03

```
In [42]: # Train the model with the base layers frozen
initial_epochs = 10
model_resnet.set_weights(initial_weights)
history_resnet = model_resnet.fit(train_image_array, train_label_array,
                                  validation_data=(valid_image_array, valid_label_array),
                                  epochs=initial_epochs)
```

Epoch 1/10  
92/92 ————— 40s 438ms/step - accuracy: 0.2098 - loss: 3.8667 - val\_accuracy: 0.6889 - val\_loss: 1.7264  
Epoch 2/10  
92/92 ————— 40s 438ms/step - accuracy: 0.5103 - loss: 1.6517 - val\_accuracy: 0.7867 - val\_loss: 0.8276  
Epoch 3/10  
92/92 ————— 40s 438ms/step - accuracy: 0.5683 - loss: 1.4143 - val\_accuracy: 0.8125 - val\_loss: 0.6622  
Epoch 4/10  
92/92 ————— 41s 446ms/step - accuracy: 0.5948 - loss: 1.3091 - val\_accuracy: 0.8410 - val\_loss: 0.5467  
Epoch 5/10  
92/92 ————— 41s 447ms/step - accuracy: 0.6039 - loss: 1.2261 - val\_accuracy: 0.8370 - val\_loss: 0.5180  
Epoch 6/10  
92/92 ————— 40s 434ms/step - accuracy: 0.6373 - loss: 1.1347 - val\_accuracy: 0.8397 - val\_loss: 0.5309  
Epoch 7/10  
92/92 ————— 41s 444ms/step - accuracy: 0.6220 - loss: 1.1613 - val\_accuracy: 0.8492 - val\_loss: 0.5059  
Epoch 8/10  
92/92 ————— 41s 444ms/step - accuracy: 0.6413 - loss: 1.1203 - val\_accuracy: 0.8302 - val\_loss: 0.5293  
Epoch 9/10  
92/92 ————— 39s 427ms/step - accuracy: 0.6550 - loss: 1.0809 - val\_accuracy: 0.8505 - val\_loss: 0.4990  
Epoch 10/10  
92/92 ————— 39s 430ms/step - accuracy: 0.6436 - loss: 1.1059 - val\_accuracy: 0.8438 - val\_loss: 0.5118

```
In [43]: def plot_performance(history, learning_rate=None, batch_size=None, finetune_epochs=None):
plt.figure(figsize=(10,5))

# history_data = history.history

# Determine whether history is keras history or a dictionary to appropriately extract the history data
if isinstance(history, keras.callbacks.History):
    history_data = history.history # Extract the history dictionary
else:
    history_data = history # Assume it's already a dictionary

# Accuracy of model training and validation vs training epoch
plt.subplot(1,2,1)
ylim_acc = [0, max(max(history_data['accuracy']),max(history_data['val_accuracy']))]
plt.plot(history_data['accuracy'], label = 'Training accuracy')
plt.plot(history_data['val_accuracy'], label = 'Validation accuracy')
plt.ylim(ylim_acc)
if finetune_epochs:
    plt.plot([finetune_epochs-1, finetune_epochs-1],plt.ylim(), label = 'Fine tuning')
else:
    pass

if learning_rate and batch_size:
    plt.title(f'Model accuracy \n lr = {learning_rate}, batch size = {batch_size}')
else: plt.title('Model accuracy')

plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(loc='lower right')

# Loss during model training and validation
plt.subplot(1,2,2)
ylim_loss = [0, max(max(history_data['loss']),max(history_data['val_loss']))]
# print(len(history_data['loss']))
plt.plot(history_data['loss'], label = 'Training loss')
plt.plot(history_data['val_loss'], label = 'Validation loss')
plt.ylim(ylim_loss)
if finetune_epochs:
    plt.plot([finetune_epochs-1, finetune_epochs-1],plt.ylim(), label = 'Fine tuning')
else:
    pass


if learning_rate and batch_size:
    plt.title(f'Model loss \n lr = {learning_rate}, batch size = {batch_size}')
else: plt.title('Model loss')
plt.ylabel('Loss')
```



```
plt.xlabel('Epoch')
plt.legend(loc='lower right')
plt.show()

print(f"The model has a training accuracy of {history_data['accuracy'][-1]*100:.2f}%\n"
      f"The model has a validation accuracy of {history_data['val_accuracy'][-1]*100:.2f}%")
return
```

In [44]: `plot_performance(history_resnet)`



The model has a training accuracy of 64.27%  
The model has a validation accuracy of 84.38%

## Model Evaluation


In [14]: `test_loss, test_acc = model_resnet.evaluate(test_image_array, test_label_array)`  
`print(f"Test accuracy: {test_acc}\n"`  
 `f"Test loss: {test_loss}")`

115/115 ————— 41s 353ms/step - accuracy: 0.8212 - loss: 0.5542  
Test accuracy: 0.8135731816291809  
Test loss: 0.5614109039306641

In [15]: `from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay`  
`prediction_array = np.argmax(model_resnet.predict(test_image_array), axis=1)`

115/115 ————— 42s 355ms/step

In [34]: `cm = confusion_matrix(test_label_array, prediction_array)`  
`disp = ConfusionMatrixDisplay(confusion_matrix=cm)`  
`disp.plot(cmap=plt.cm.Blues) # You can change the color map as desired`  
`fig = disp.ax_.get_figure()`  
`fig.set_figwidth(12); fig.set_figheight(10)`  
`plt.title("Confusion Matrix - ResNet")`  
`plt.xticks(rotation=90, ha='right') # Rotate x labels for better readability`  
`plt.yticks(rotation=0) # Keep y labels horizontal`  
`plt.tight_layout() # Adjust layout to make room for rotated labels`  
`plt.show()`  
  
`label_dict`





```
Out[34]: {0: 'Abyssinian',
1: 'american_bulldog',
2: 'american_pit_bull_terrier',
3: 'basset_hound',
4: 'beagle',
5: 'Bengal',
6: 'Birman',
7: 'Bombay',
8: 'boxer',
9: 'British_Shorthair',
10: 'chihuahua',
11: 'Egyptian_Mau',
12: 'english_cocker_spaniel',
13: 'english_setter',
14: 'german_shorthaired',
15: 'great_pyrenees',
16: 'havanese',
17: 'japanese_chin',
18: 'keeshond',
19: 'leonberger',
20: 'Maine_Coon',
21: 'miniature_pinscher',
22: 'newfoundland',
23: 'Persian',
24: 'pomeranian',
25: 'pug',
26: 'Ragdoll',
27: 'Russian_Blue',
28: 'saint_bernard',
29: 'samoyed',
30: 'scottish_terrier',
31: 'shiba_inu',
32: 'Siamese',
33: 'Sphynx',
34: 'staffordshire_bull_terrier',
35: 'wheaten_terrier',
36: 'yorkshire_terrier'}
```

## Model Prediction and Visualization

```
In [ ]: # Sample random images and their indices
num_samples = 9 # number of samples
num_rows = int(round(sqrt(num_samples))); num_cols = int(num_samples/num_rows) # number of rows and column:
rand = random.randint(num_test,size = (num_samples)) # random index for c

image_test_rand_array = test_image_array[rand]
label_test_rand_array = test_label_array[rand]
prediction_rand_array = np.argmax(model_resnet.predict(image_test_rand_array),axis=1)


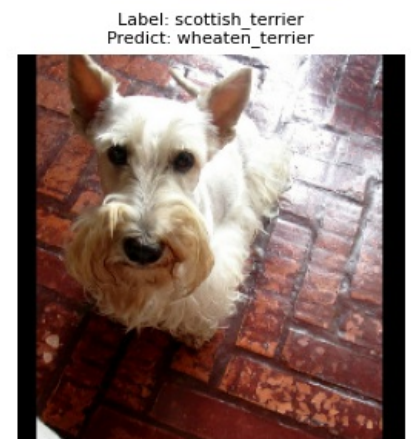


plt.figure(figsize=(num_rows*3,num_cols*3))
# fig, axes1 = plt.subplots(num_rows,num_cols,figsize=(num_rows*2,num_cols*2))

for i in range(num_rows):
    for j in range(num_cols):
        index = i * num_cols + j
        plt.subplot(num_rows,num_cols,index+1)
        image = image_test_rand_array[index]/255.0 # Extract the image
        label = label_test_rand_array[index] # Extract the label
        prediction = prediction_rand_array[index]

        # Original pictures (no augmentation layer applied)
        plt.axis("off")
        # Display the image
        plt.imshow(image)
        plt.title(f"Label: {label_dict[label]}\n"
                  f"Predict: {label_dict[prediction]}",
                  fontsize = 8)

plt.tight_layout()
```

1/1 ————— 0s 128ms/step



## InceptionV3-based model

```
In [20]: import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.applications import InceptionV3
from tensorflow.keras.applications.inception_v3 import preprocess_input as inception_preprocess
```

Since the inception model use similar input image shapes as ResNet, I will reuse the processed train, validation, and test sets

## Base Model Construction

```
In [21]: base_model = InceptionV3(weights='imagenet',
                                include_top=False)

# Freeze the weights of the model
```

```
base_model.trainable = False
```

```
## Investigate the structure of the base model and make sure that the weights are frozen  
# base_model.summary()
```

## Full Model

```
In [24]: model_inception = models.Sequential([  
    layers.Input(TARGET_SIZE + (3,)),  
    layers.Lambda(inception_preprocess),  
    base_model,  
    layers.GlobalAveragePooling2D(),  
    layers.BatchNormalization(),  
    layers.Dropout(0.5),  
    layers.Dense(num_classes, activation='relu'),  
    layers.Dropout(0.5),  
    layers.Dense(num_classes, activation='sigmoid') # Use 1 for binary classification (dogs vs cats)  
)  
  
initial_weight_inception = model_inception.get_weights()  
  
model_inception.summary()
```

Model: "sequential\_2"

| Layer (type)   | Output Shape        | Param #    |
|--|---------------------|------------|
| lambda_2 (Lambda)                                      | (None, 224, 224, 3) | 0          |
| inception_v3 (Functional)                              | (None, 5, 5, 2048)  | 21,802,784 |
| global_average_pooling2d_2<br>(GlobalAveragePooling2D) | (None, 2048)        | 0          |
| batch_normalization_96<br>(BatchNormalization)         | (None, 2048)        | 8,192      |
| dropout_4 (Dropout)                                    | (None, 2048)        | 0          |
| dense_4 (Dense)  | (None, 37)          | 75,813     |
| dropout_5 (Dropout)                                    | (None, 37)          | 0          |
| dense_5 (Dense)  | (None, 37)          | 1,406      |

Total params: 21,888,195 (83.50 MB)

Trainable params: 81,315 (317.64 KB)

Non-trainable params: 21,806,880 (83.19 MB)

## Model Training

```
In [25]: model_inception.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.005),  
    loss=tf.keras.losses.SparseCategoricalCrossentropy,  
    metrics=["accuracy"])  
loss0, acc0 = model_inception.evaluate(test_image_array, test_label_array)  
  
print("initial loss: {:.2f}".format(loss0))  
print("initial accuracy: {:.2f}".format(acc0))
```

115/115 ————— 22s 180ms/step - accuracy: 0.0130 - loss: 3.9160  
initial loss: 3.93  
initial accuracy: 0.01


```
In [26]: # Train the model with the base layers frozen  
initial_epochs = 10  
model_inception.set_weights(initial_weight_inception)  
history_inception = model_inception.fit(train_image_array, train_label_array,  
    validation_data=(valid_image_array, valid_label_array),  
    epochs=initial_epochs)
```

```

Epoch 1/10
92/92 ————— 25s 236ms/step - accuracy: 0.3012 - loss: 3.2835 - val_accuracy: 0.7364 - val_loss: 1.1180
Epoch 2/10
92/92 ————— 22s 236ms/step - accuracy: 0.5874 - loss: 1.4721 - val_accuracy: 0.7948 - val_loss: 0.6620
Epoch 3/10
92/92 ————— 22s 236ms/step - accuracy: 0.6440 - loss: 1.2095 - val_accuracy: 0.8274 - val_loss: 0.6148
Epoch 4/10
92/92 ————— 22s 240ms/step - accuracy: 0.6618 - loss: 1.0902 - val_accuracy: 0.8370 - val_loss: 0.5594
Epoch 5/10
92/92 ————— 22s 234ms/step - accuracy: 0.6772 - loss: 1.0283 - val_accuracy: 0.8505 - val_loss: 0.5497
Epoch 6/10
92/92 ————— 21s 233ms/step - accuracy: 0.6888 - loss: 1.0056 - val_accuracy: 0.8682 - val_loss: 0.5698
Epoch 7/10
92/92 ————— 21s 232ms/step - accuracy: 0.6882 - loss: 1.0169 - val_accuracy: 0.8478 - val_loss: 0.5656
Epoch 8/10
92/92 ————— 21s 233ms/step - accuracy: 0.7027 - loss: 0.9206 - val_accuracy: 0.8315 - val_loss: 0.6064
Epoch 9/10
92/92 ————— 21s 232ms/step - accuracy: 0.6878 - loss: 0.9712 - val_accuracy: 0.8315 - val_loss: 0.5811
Epoch 10/10
92/92 ————— 22s 235ms/step - accuracy: 0.6820 - loss: 1.0087 - val_accuracy: 0.8166 - val_loss: 0.5891

```

```
In [27]: plot_performance(history_inception)
```



The model has a training accuracy of 69.02%  
The model has a validation accuracy of 81.66%

## Model Evaluation

```
In [30]: test_loss_inception, test_acc_inception = model_inception.evaluate(test_image_array, test_label_array)
print(f"Test accuracy: {test_acc}\n"
      f"Test loss: {test_loss}")

prediction_array_inception = np.argmax(model_inception.predict(test_image_array), axis=1)
```

```

115/115 ————— 21s 180ms/step - accuracy: 0.8343 - loss: 0.4823
Test accuracy: 0.8135731816291809
Test loss: 0.5614109039306641
115/115 ————— 23s 189ms/step


```

```
In [35]: cm = confusion_matrix(test_label_array, prediction_array_inception)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap=plt.cm.Blues) # You can change the color map as desired
fig = disp.ax_.get_figure()
fig.set_figwidth(12)
fig.set_figheight(10)
```



```
plt.title("Confusion Matrix - Inception model")
plt.xticks(rotation=90, ha='right') # Rotate x labels for better readability
plt.yticks(rotation=0) # Keep y labels horizontal
plt.tight_layout() # Adjust layout to make room for rotated labels
plt.show()
```

label\_dict





```
Out[35]: {0: 'Abyssinian',
1: 'american_bulldog',
2: 'american_pit_bull_terrier',
3: 'basset_hound',
4: 'beagle',
5: 'Bengal',
6: 'Birman',
7: 'Bombay',
8: 'boxer',
9: 'British_Shorthair',
10: 'chihuahua',
11: 'Egyptian_Mau',
12: 'english_cocker_spaniel',
13: 'english_setter',
14: 'german_shorthaired',
15: 'great_pyrenees',
16: 'havanese',
17: 'japanese_chin',
18: 'keeshond',
19: 'leonberger',
20: 'Maine_Coon',
21: 'miniature_pinscher',
22: 'newfoundland',
23: 'Persian',
24: 'pomeranian',
25: 'pug',
26: 'Ragdoll',
27: 'Russian_Blue',

28: 'saint_bernard',
29: 'samoyed',
30: 'scottish_terrier',
31: 'shiba_inu',
32: 'Siamese',
33: 'Sphynx',
34: 'staffordshire_bull_terrier',
35: 'wheaten_terrier',
36: 'yorkshire_terrier'}
```

## Model Prediction and Visualization

```
In [36]: # Sample random images and their indices
num_samples = 9 # number of samples
num_rows = int(round(sqrt(num_samples))); num_cols = int(num_samples/num_rows) # number of rows and column:
rand = random.randint(num_test,size = (num_samples)) # random index for c

image_test_rand_array = test_image_array[rand]
label_test_rand_array = test_label_array[rand]
prediction_rand_array = np.argmax(model_inception.predict(image_test_rand_array),axis=1)

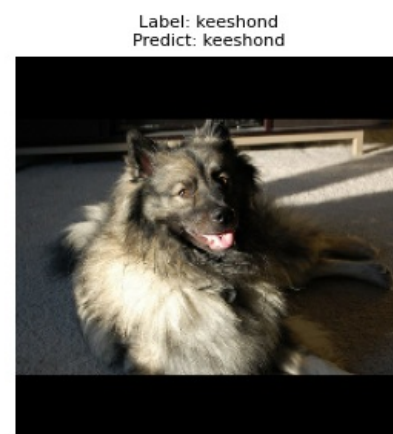
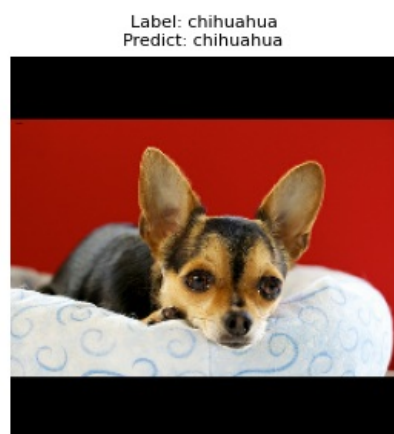
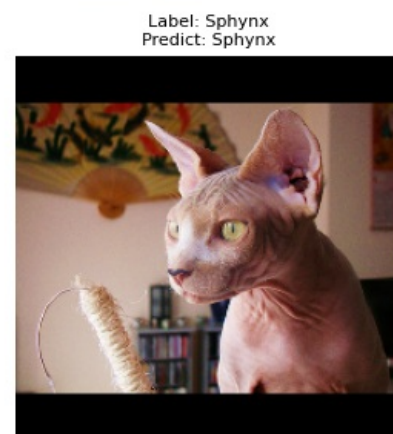

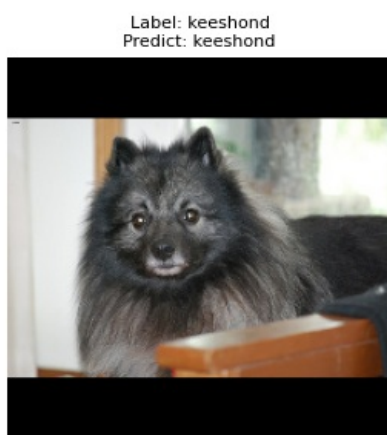
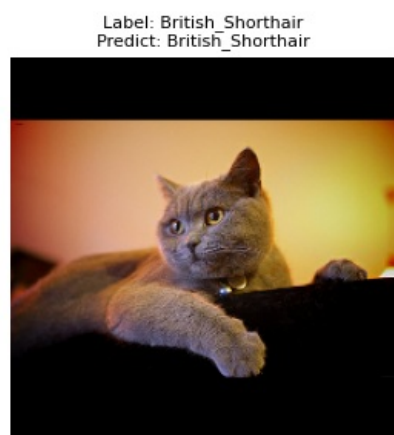
plt.figure(figsize=(num_rows*3,num_cols*3))

for i in range(num_rows):
    for j in range(num_cols):
        index = i * num_cols + j
        plt.subplot(num_rows,num_cols,index+1)
        image = image_test_rand_array[index]/255.0 # Extract the image
        label = label_test_rand_array[index] # Extract the label
        prediction = prediction_rand_array[index]

        # Original pictures (no augmentation layer applied)
        plt.axis("off")
        # Display the image
        plt.imshow(image)
        plt.title(f"Label: {label_dict[label]}\n"
                  f"Predict: {label_dict[prediction]}",
                  fontsize = 8)

plt.tight_layout()
```

1/1 ————— 0s 87ms/step



## ResNet vs Inception Comparison

### Model Performance Comparison

```
In [50]: test_loss_resnet, test_acc_resnet = model_resnet.evaluate(test_image_array, test_label_array)
         test_loss_inception, test_acc_inception = model_inception.evaluate(test_image_array, test_label_array)
```

```
115/115 ————— 40s 346ms/step - accuracy: 0.8325 - loss: 0.4967
115/115 ————— 22s 191ms/step - accuracy: 0.8343 - loss: 0.4823
```

```
In [56]: msg_loss = "comparable"; msg_acc = "comparable"
         test_loss_diff = test_loss_resnet - test_loss_inception
         test_acc_diff = test_acc_resnet - test_acc_inception
         diff_threshold = 0.01

         if np.abs(test_loss_diff) > diff_threshold:
```

```

    if test_loss_resnet > test_loss_inception: msg_loss = "better"
    else: msg_loss = "worse"

if np.abs(test_acc_diff) > diff_threshold:
    if test_acc_resnet > test_acc_inception: msg_acc = "better"
    else: msg_acc = "worse"

print(f"The ResNet-based model has {msg_acc} accuracy compared to the Inception-based model\n"
      f"Resnet-based model accuracy: {test_acc_resnet*100:.2f}%\n"
      f"Inception-based model accuracy: {test_acc_inception*100:.2f}%")

print(f"The ResNet-based model has {msg_loss} accuracy compared to the Inception-based model\n"
      f"Resnet-based model loss: {test_loss_resnet*100:.2f}%\n"
      f"Inception-based model loss: {test_loss_inception*100:.2f}%")

```

The ResNet-based model has comparable accuracy compared to the Inception-based model  
 Resnet-based model accuracy: 83.32%  
 Inception-based model accuracy: 83.29%  
 The ResNet-based model has comparable accuracy compared to the Inception-based model  
 Resnet-based model loss: 50.15%  
 Inception-based model loss: 50.80%

## Prediction Comparison

```

In [57]: # Sample random images and their indices
num_samples = 25                                     # number of sample:
num_rows = int(round(sqrt(num_samples))); num_cols = int(num_samples/num_rows) # number of rows and column:
rand = random.randint(num_test,size = (num_samples)) # random index for ci

image_test_rand_array = test_image_array[rand]
label_test_rand_array = test_label_array[rand]
prediction_rand_array_resnet = np.argmax(model_resnet.predict(image_test_rand_array),axis=1)
prediction_rand_array_inception = np.argmax(model_inception.predict(image_test_rand_array),axis=1)

plt.figure(figsize=(num_rows*3,num_cols*3))
# fig, axes1 = plt.subplots(num_rows,num_cols,figsize=(num_rows*2,num_cols*2))

for i in range(num_rows):
    for j in range(num_cols):
        index = i * num_cols + j
        plt.subplot(num_rows,num_cols,index+1)
        image = image_test_rand_array[index]/255.0 # Extract the image
        label = label_test_rand_array[index] # Extract the label
        prediction_resnet = prediction_rand_array_resnet[index]
        prediction_inception = prediction_rand_array_inception[index]

        # Original pictures (no augmentation layer applied)
        plt.axis("off")
        # Display the image
        plt.imshow(image)
        plt.title(f"Label: {label_dict[label]}\n"
                  f"ResNet predicts: {label_dict[prediction_resnet]} \n"
                  f"Inception predicts: {label_dict[prediction_inception]}",
                  fontsize = 8)


plt.tight_layout()

```


1/1 ————— 0s 287ms/step  
 1/1 ————— 0s 160ms/step




Label: great\_pyrenees  
ResNet predicts: great\_pyrenees  
Inception predicts: great\_pyrenees



Label: basset\_hound  
ResNet predicts: basset\_hound  
Inception predicts: basset\_hound




Label: american\_pit\_bull\_terrier  
ResNet predicts: american\_pit\_bull\_terrier  
Inception predicts: american\_pit\_bull\_terrier



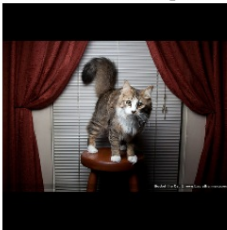
Label: Egyptian\_Mau  
ResNet predicts: Egyptian\_Mau  
Inception predicts: Egyptian\_Mau




Label: leonberger  
ResNet predicts: leonberger  
Inception predicts: wheaten\_terrier



Label: Maine\_Coon  
ResNet predicts: Maine\_Coon  
Inception predicts: wheaten\_terrier




Label: japanese\_chin  
ResNet predicts: japanese\_chin  
Inception predicts: japanese\_chin




Label: wheaten\_terrier  
ResNet predicts: wheaten\_terrier  
Inception predicts: wheaten\_terrier




Label: Bombay  
ResNet predicts: Abyssinian  
Inception predicts: pomeranian



Label: yorkshire\_terrier  
ResNet predicts: yorkshire\_terrier  
Inception predicts: havanese



Label: wheaten\_terrier  
ResNet predicts: wheaten\_terrier  
Inception predicts: wheaten\_terrier




Label: Persian  
ResNet predicts: Persian  
Inception predicts: Persian



Label: american\_bulldog  
ResNet predicts: american\_bulldog  
Inception predicts: american\_bulldog




Label: scottish\_terrier  
ResNet predicts: scottish\_terrier  
Inception predicts: scottish\_terrier




Label: Bombay  
ResNet predicts: Bombay  
Inception predicts: Bombay




Label: english\_cocker\_spaniel  
ResNet predicts: english\_cocker\_spaniel  
Inception predicts: english\_cocker\_spaniel



Label: leonberger  
ResNet predicts: leonberger  
Inception predicts: leonberger



Label: shiba\_inu  
ResNet predicts: shiba\_inu  
Inception predicts: shiba\_inu




Label: Abyssinian  
ResNet predicts: Abyssinian  
Inception predicts: Abyssinian




Label: pomeranian  
ResNet predicts: pomeranian  
Inception predicts: pomeranian




Label: leonberger  
ResNet predicts: leonberger  
Inception predicts: leonberger




Label: havanese  
ResNet predicts: havanese  
Inception predicts: havanese



Label: english\_setter  
ResNet predicts: english\_cocker\_spaniel  
Inception predicts: english\_setter



Label: Siamese  
ResNet predicts: Siamese  
Inception predicts: Siamese



Label: scottish\_terrier  
ResNet predicts: wheaten\_terrier  
Inception predicts: scottish\_terrier

